

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load data

Amazon_sales_data = pd.read_csv(r"C:\Users\steph\OneDrive\Documents\
Desktop\Amazon sales data\amazon_sales_data 2025.csv")

# Show first few rows

Amazon_sales_data.head()
```

	Order ID	Date	Product	Category	Price	Quantity	\
0	ORD0001	14-03-2025	Running Shoes	Footwear	60	3	
1	ORD0002	20-03-2025	Headphones	Electronics	100	4	
2	ORD0003	15-02-2025	Running Shoes	Footwear	60	2	
3	ORD0004	19-02-2025	Running Shoes	Footwear	60	3	
4	ORD0005	10-03-2025	Smartwatch	Electronics	150	3	

	Total Sales	Customer Name	Customer Location	Payment Method	Status
0	180	Emma Clark	New York	Debit Card	Cancelled
1	400	Emily Johnson	San Francisco	Debit Card	Pending
2	120	John Doe	Denver	Amazon Pay	Cancelled
3	180	Olivia Wilson	Dallas	Credit Card	Pending
4	450	Emma Clark	New York	Debit Card	Pending

# Dataset Summary

```
Amazon_sales_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Order ID              250 non-null   object
1   Date                  250 non-null   object
2   Product               250 non-null   object
3   Category              250 non-null   object
4   Price                 250 non-null   int64
5   Quantity              250 non-null   int64
6   Total Sales           250 non-null   int64
7   Customer Name         250 non-null   object
8   Customer Location     250 non-null   object
```

```

9    Payment Method      250 non-null    object
10   Status              250 non-null    object
dtypes: int64(3), object(8)
memory usage: 21.6+ KB

```

*## Checking the correct Date format*

```

Amazon_sales_data["Date"] =
pd.to_datetime(Amazon_sales_data["Date"], format="%d-%m-%Y")

```

*## Rename Columns for Clarity*

```

Amazon_sales_data.rename(columns={Amazon_sales_data.columns[1]: "Order_
Date"}, inplace=True)

```

```

Amazon_sales_data["Order_Date"] =
pd.to_datetime(Amazon_sales_data["Order_Date"], errors = "coerce")

```

*# Converts string to datetime object. errors="coerce" replaces invalid formats with NaT (null),  
# which is then dropped in step 1.*

```

Amazon_sales_data["Order_Date"].head()

```

```

0    2025-03-14
1    2025-03-20
2    2025-02-15
3    2025-02-19
4    2025-03-10
Name: Order_Date, dtype: datetime64[ns]

```

*# Strip Whitespaces from All Object Columns  
# Why: Removes leading/trailing whitespace that may cause grouping errors or  
# mismatched filtering in Power BI.*

```

Amazon_sales_data=Amazon_sales_data.apply(lambda x:x.str.strip() if
x.dtype== "object" else x)

```

	Order ID	Order_Date	Product	Category	Price	
Quantity \						
0	ORD0001	2025-03-14	Running Shoes	Footwear	60	3
1	ORD0002	2025-03-20	Headphones	Electronics	100	4
2	ORD0003	2025-02-15	Running Shoes	Footwear	60	2
3	ORD0004	2025-02-19	Running Shoes	Footwear	60	3
4	ORD0005	2025-03-10	Smartwatch	Electronics	150	3

...	...	...	...	...	...	...
245	ORD0246	2025-03-17	T-Shirt	Clothing	20	2
246	ORD0247	2025-03-30	Jeans	Clothing	40	1
247	ORD0248	2025-03-05	T-Shirt	Clothing	20	2
248	ORD0249	2025-03-08	Smartwatch	Electronics	150	3
249	ORD0250	2025-02-19	Smartphone	Electronics	500	4

	Total Sales	Customer Name	Customer Location	Payment Method	Status
0	180	Emma Clark	New York	Debit Card	Cancelled
1	400	Emily Johnson	San Francisco	Debit Card	Pending
2	120	John Doe	Denver	Amazon Pay	Cancelled
3	180	Olivia Wilson	Dallas	Credit Card	Pending
4	450	Emma Clark	New York	Debit Card	Pending
...	...	...	...	...	...
245	40	Daniel Harris	Miami	Debit Card	Cancelled
246	40	Sophia Miller	Dallas	Debit Card	Cancelled
247	40	Chris White	Denver	Debit Card	Cancelled
248	450	Emily Johnson	New York	Debit Card	Cancelled
249	2000	Emily Johnson	Seattle	Amazon Pay	Completed

[250 rows x 11 columns]

```
# Standardize Column Names
# Result: Avoids syntax errors in Power BI and Python due to
# inconsistent or
# space-containing column headers.

Amason_sales_data.columns =
Amason_sales_data.columns.str.strip().str.replace(" ", "_") # for
column name

Amason_sales_data.columns
```

```

Index(['Order_ID', 'Order_Date', 'Product', 'Category', 'Price',
      'Quantity',
      'Total_Sales', 'Customer_Name', 'Customer_Location',
      'Payment_Method',
      'Status'],
      dtype='object')

# Dropping Null/Empty Rows
# Reason: These are essential columns for time series, sales
aggregation,
# and trend visualization. Missing values here would disrupt EDA and
Power BI integration.

Amazon_sales_data.dropna(subset=['Order_Date', 'Price', 'Quantity',
'Total_Sales'], inplace=True)

Amazon_sales_data.to_csv("cleaned_amazon_sales.csv", index=False) #
this cleaned data for PowerBI Analysis

## Exploratory Data Analysis (EDA)

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

## 1. best performing product categories

best_product_categories = Amazon_sales_data.groupby("Category")
[["Total_Sales"]].sum().sort_values(by =
"Total_Sales", ascending=False)

best_product_categories

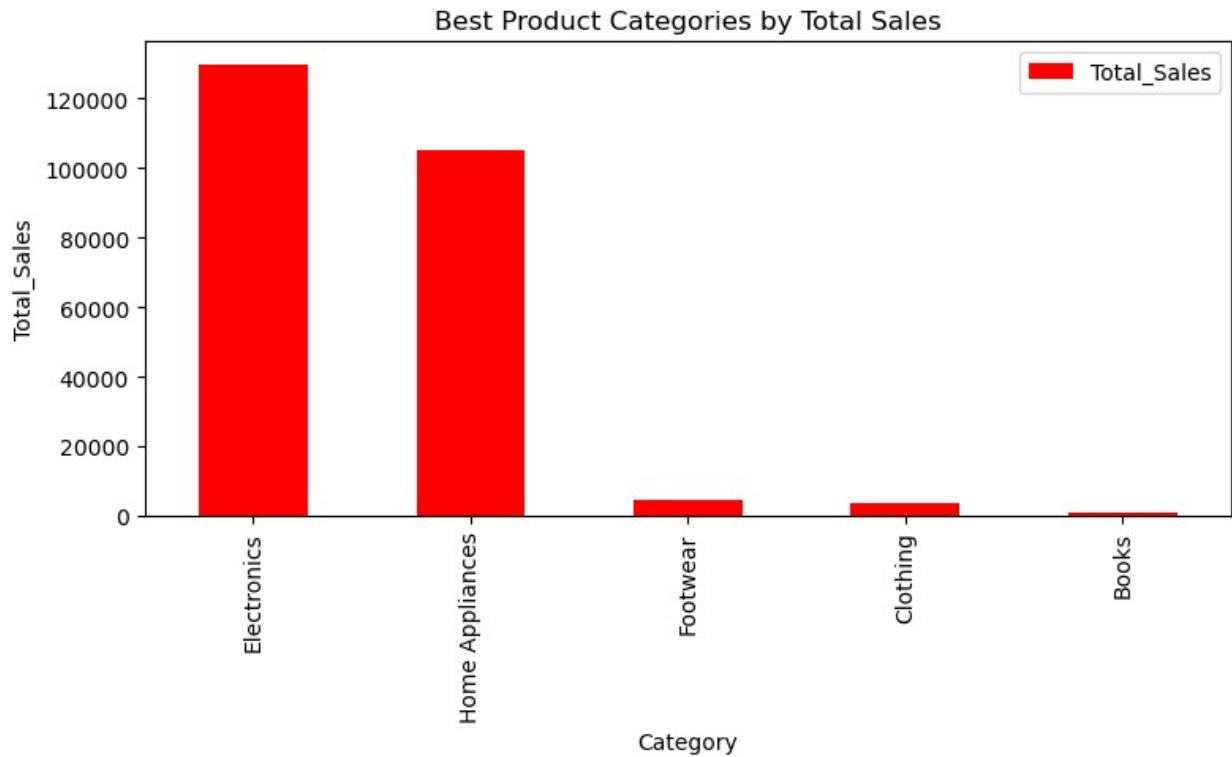

```

Category	Total_Sales
Electronics	129950
Home Appliances	105000
Footwear	4320
Clothing	3540
Books	1035

```

best_product_categories.plot(kind = "bar", color = "Red", figsize=(8,5))
plt.xlabel("Category")
plt.ylabel("Total_Sales")
plt.title("Best Product Categories by Total Sales")
plt.tight_layout()
plt.show()

```



## ## 2.Sales Trend Over Time

```
daily_sales = Amason_sales_data.groupby("Order_Date")  
["Total_Sales"].sum().sort_index()
```


daily\_sales

Order_Date	
2025-02-02	3600
2025-02-03	3360
2025-02-04	6815
2025-02-05	5400
2025-02-06	11400
2025-02-07	2520
2025-02-08	1640
2025-02-09	3550
2025-02-10	10965
2025-02-11	3550
2025-02-12	3560
2025-02-13	4860
2025-02-14	1015
2025-02-15	520
2025-02-16	9540
2025-02-17	1005
2025-02-18	8810
2025-02-19	2195
2025-02-20	6730

2025-02-21	4600
2025-02-22	1800
2025-02-23	2570
2025-02-24	6900
2025-02-25	6300
2025-02-26	1980
2025-02-27	300
2025-02-28	7210
2025-03-01	1800
2025-03-02	4575
2025-03-03	120
2025-03-04	5900
2025-03-05	2900
2025-03-06	9720
2025-03-07	8660
2025-03-08	3125
2025-03-09	200
2025-03-10	3965
2025-03-11	1240
2025-03-12	710
2025-03-13	4160
2025-03-14	4350
2025-03-15	7190
2025-03-16	2735
2025-03-17	540
2025-03-18	3700
2025-03-19	4115
2025-03-20	3630
2025-03-21	1960
2025-03-22	950
2025-03-23	3900
2025-03-24	9520
2025-03-25	6015
2025-03-26	2970
2025-03-27	15
2025-03-28	2400
2025-03-29	2600
2025-03-30	7465
2025-03-31	6600
2025-04-01	320
2025-04-02	3100

Name: Total\_Sales, dtype: int64

```
daily_sales.plot(kind = "line",color = "Green",figsize= (8,5))
plt.xlabel("Order_Date")
plt.ylabel("Total_Sales")
plt.title("Sales Trend Over Time")
plt.tight_layout()
plt.show()
```



### ## 3.Sales By location

```
Amason_sales_data.columns
```

```
Index(['Order_ID', 'Order_Date', 'Product', 'Category', 'Price',
      'Quantity',
      'Total_Sales', 'Customer_Name', 'Customer_Location',
      'Payment_Method',
      'Status'],
      dtype='object')
```


```
sales_location = Amason_sales_data.groupby("Customer_Location")
[["Total_Sales"]].sum().sort_values("Total_Sales",ascending=False)
```

```
sales_location
```

Customer_Location	Total_Sales
Miami	31700
Denver	29785
Houston	28390
Dallas	27145
Seattle	26890
Boston	26170
Chicago	20810
New York	18940

Los Angeles	17820
San Francisco	16195

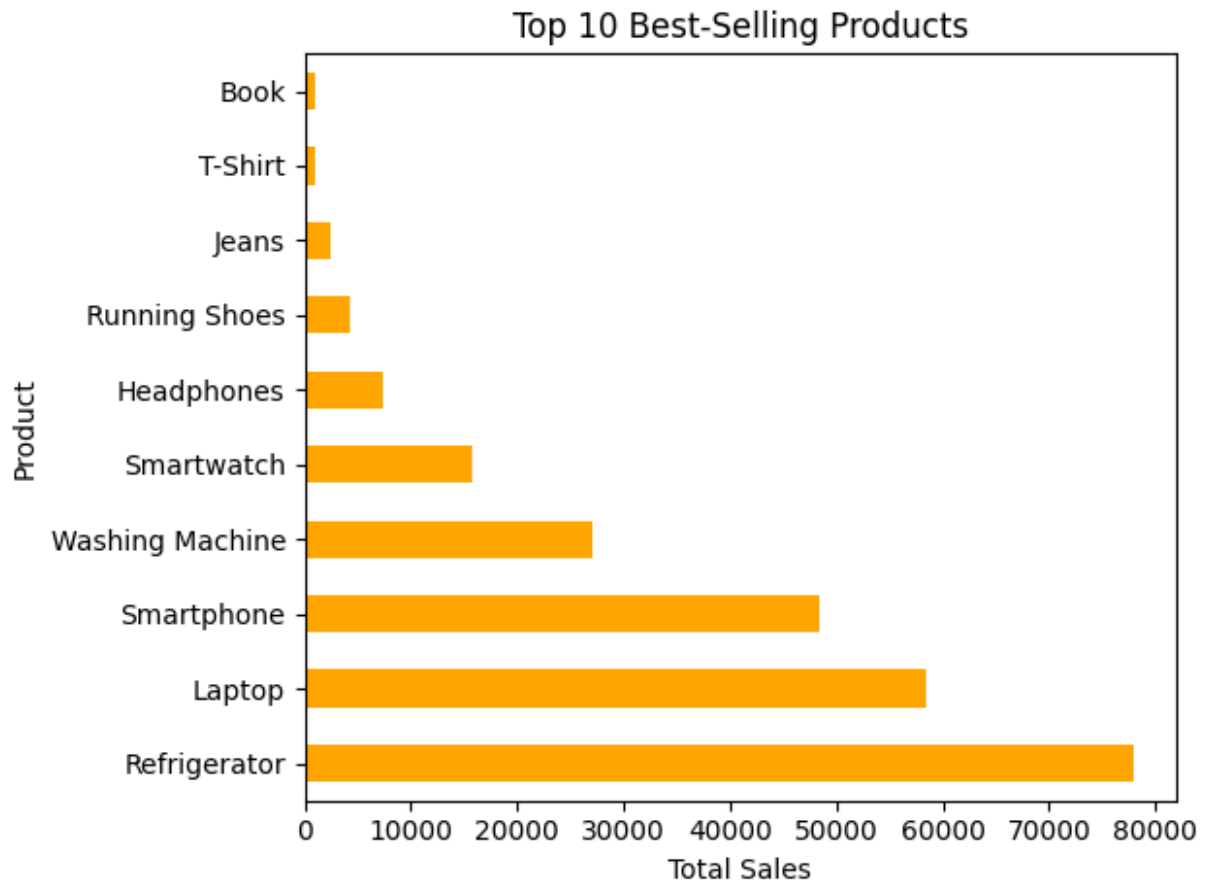
```
sales_location.plot(kind = "bar",color = "salmon",figsize=(8,5))
plt.xlabel("Customer_location")
plt.title("Sales BY Customer Location")
plt.ylabel("Total_Sales")
plt.tight_layout()
plt.show()
```



#### ## 4.Top Products by Sales:

```
top_products = Amason_sales_data.groupby("Product")
["Total_Sales"].sum().sort_values(ascending=False).head(10)
top_products.plot(kind="barh", color="orange")
plt.title("Top 10 Best-Selling Products")
plt.xlabel("Total Sales")
plt.tight_layout()
plt.show()
```






*# 5.Payment Method Preference:*

```
Amason_sales_data["Payment_Method"].value_counts().plot(kind="pie",  
autopct='%1.1f%%', figsize=(6,6))  
plt.title("Payment Method Distribution")  
plt.ylabel("")  
plt.show()
```

Payment Method Distribution



# 6. Order Status Breakdown:

```
import seaborn as sns
```

```
sns.countplot(data=Amason_sales_data, x="Status", hue="Status",  
palette="Set2", legend=False)  
plt.title("Order Status Count")  
plt.show()
```



### *## 7.Order Status Distribution*

```
Amazon_sales_data.columns
```

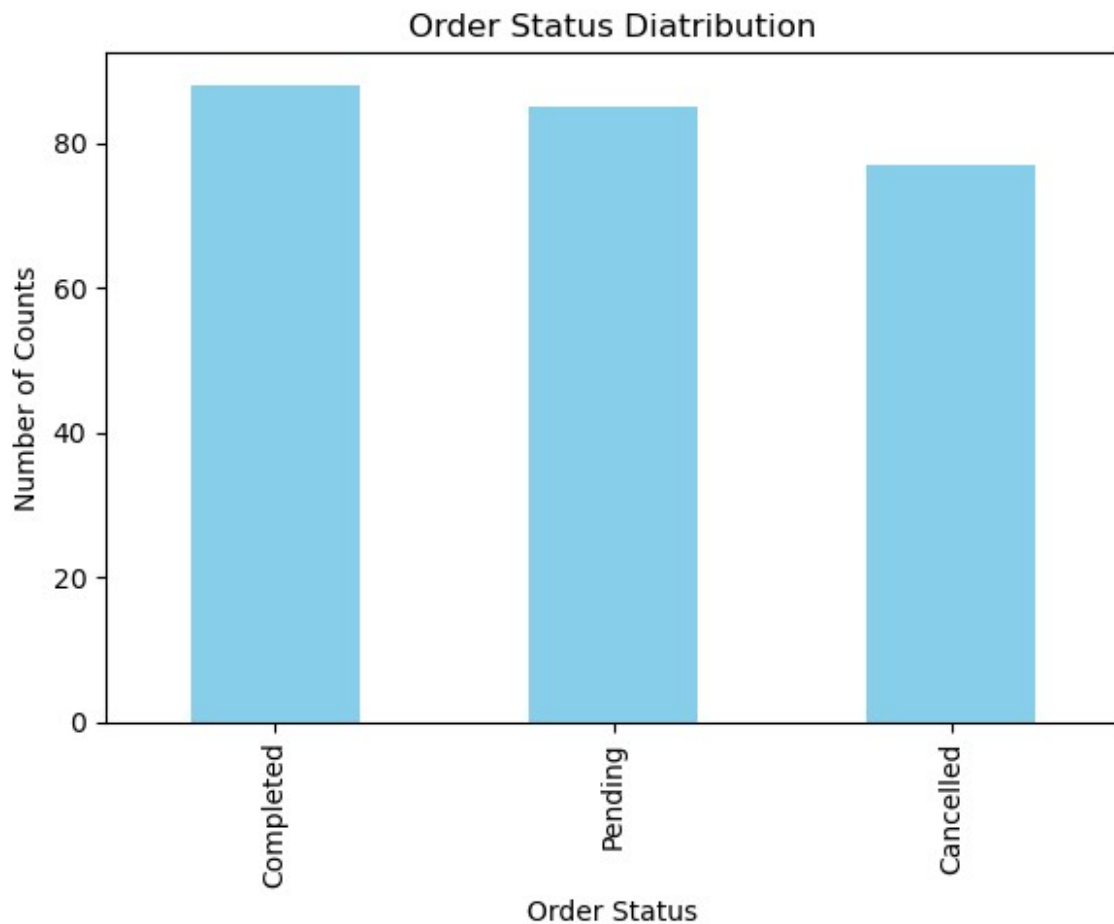
```
Index(['Order_ID', 'Order_Date', 'Product', 'Category', 'Price',  
      'Quantity',  
      'Total_Sales', 'Customer_Name', 'Customer_Location',  
      'Payment_Method',  
      'Status'],  
      dtype='object')
```

```
Amazon_sales_data["Status"].value_counts()
```

```
Status  
Completed      88  
Pending        85  
Cancelled      77  
Name: count, dtype: int64
```

```
Amazon_sales_data["Status"].value_counts().plot(kind =  
"bar",figsize=(6,5),color ="skyblue")  
plt.xlabel("Order Status")  
plt.ylabel("Number of Counts")  
plt.title("Order Status Distribution")
```

```
plt.tight_layout()
plt.show()
```



```
### Statistical Insights with scipy and statsmodels
```

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
### correlation Matrix
```

```
cor = Amason_sales_data[["Price", "Quantity", "Total_Sales"]].corr()
```

```
cor
```

	Price	Quantity	Total_Sales
Price	1.000000	-0.010858	0.846673
Quantity	-0.010858	1.000000	0.332444
Total_Sales	0.846673	0.332444	1.000000


+1 = Perfect Positive Correlation, As price increases, total sales also tend to increase. logical, because  $\text{Total\_Sales} = \text{Price} \times \text{Quantity}$

0 = No Correlation

-1 = Perfect Negative Correlation

Price and Quantity aren't related A product's price doesn't influence how many are bought in this dataset

```
#In this data, Total Sales depends more on Price than Quantity, and Price doesn't affect Quantity bought.  
# by heat map we can see clearly  
sns.heatmap(cor,annot=True,cmap="coolwarm") # annot=True This means  
annotate the cells – display the numeric values (correlation  
coefficients) inside each box.  
plt.title("Correlation Heatmap")  
plt.show()
```



*# ANOVA: Are Total Sales significantly different between categories?*

```
from scipy import stats
```

*# Grouping sales by category*

```
Amazon_sales_data.columns
```

```
Index(['Order_ID', 'Order_Date', 'Product', 'Category', 'Price',
       'Quantity',
       'Total_Sales', 'Customer_Name', 'Customer_Location',
       'Payment_Method',
       'Status'],
      dtype='object')
```

```
sales_category = [group["Total_Sales"].values for name, group in
Amazon_sales_data.groupby("Category")]
```

```
sales_category
```

```
[array([30, 15, 75, 15, 15, 15, 75, 45, 15, 75, 75, 75, 45, 15, 60,
        15, 15,
        15, 75, 75, 15, 75, 30, 60, 15], dtype=int64),
 array([ 20,  60, 160,  20,  20, 100,  80,  40, 100, 120,  80, 100,
        200,
        40, 200,  40,  60, 120, 200,  20,  60, 200,  80,  20,  80,
```

```

40,
    200, 40, 100, 20, 160, 160, 200, 80, 80, 80, 40, 40,
40,
    40], dtype=int64),
array([ 400, 450, 600, 500, 500, 1600, 600, 1000, 400, 300,
100,
    500, 500, 2000, 2400, 300, 2400, 1600, 150, 300, 150,
2000,
    2400, 450, 300, 500, 500, 300, 4000, 2500, 300, 300,
100,
    750, 2400, 2500, 200, 200, 300, 300, 1500, 2500, 750,
750,
    3200, 1000, 600, 450, 300, 1000, 200, 2500, 3200, 3200,
500,
    1000, 4000, 500, 2400, 500, 300, 750, 1000, 800, 2500,
450,
    750, 1600, 300, 1000, 1500, 300, 300, 400, 150, 750,
200,
    600, 1500, 2400, 100, 1000, 750, 4000, 800, 300, 300,
1500,
    2400, 2400, 300, 1500, 500, 3200, 1500, 1000, 200, 2400,
500,
    800, 750, 400, 300, 450, 500, 400, 4000, 2500, 300,
100,
    750, 150, 2000, 2000, 2000, 800, 450, 2000], dtype=int64),
array([180, 120, 180, 180, 120, 240, 120, 240, 300, 300, 60, 60,
180,
    300, 180, 120, 120, 60, 120, 300, 60, 120, 240, 60, 180,
60,
    120], dtype=int64),
array([1800, 1200, 4800, 3600, 600, 1800, 2400, 3600, 1200, 4800,
4800,
    4800, 600, 2400, 2400, 6000, 3000, 1200, 6000, 2400, 2400,
2400,
    4800, 1800, 2400, 2400, 2400, 1200, 3600, 600, 2400, 3600,
3000,
    600, 3600, 1200, 1200, 1200, 3600, 1200], dtype=int64)]

f_stat,p_val = stats.f_oneway(*sales_category)      # Or anova_result
= stats.f_oneway(*category_groups)                # print("ANOVA F-
statistic:", anova_result.statistic)
                                                    # print("ANOVA p-
value:", anova_result.pvalue)

print("f_statistics:",f_stat)
print("p_value:",p_val)

f_statistics: 53.463921351737696
p_value: 2.4079237572585064e-32

```

```

if p_val<0.05:
    print("Significant difference found between categories.")
else:
    print("no Significant difference found between categories.")

```

Significant difference found between categories.

```

import statsmodels.api as sm

# Linear Regression: Predict Total Sales using Price & Quantity
# Predicting sales from price and quantity linear Regression is useful
X = Amason_sales_data[["Price","Quantity"]]
Y = Amason_sales_data["Total_Sales"]
X = sm.add_constant(X) ## Add Constant(intercept)
model = sm.OLS(Y,X).fit()
print(model.summary())

```

#### OLS Regression Results

```

=====
=====
Dep. Variable:          Total_Sales    R-squared:
0.834
Model:                  OLS           Adj. R-squared:
0.832
Method:                 Least Squares   F-statistic:
618.6
Date:                   Fri, 30 May 2025   Prob (F-statistic):
6.56e-97
Time:                   21:43:36          Log-Likelihood:
-1913.2
No. Observations:      250              AIC:
3832.
Df Residuals:          247              BIC:
3843.
Df Model:               2

Covariance Type:       nonrobust

=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					
-----					
-----					



const	-840.4835	78.556	-10.699	0.000	-995.208
-685.759					
Price	2.7974	0.085	32.760	0.000	2.629
2.966					
Quantity	299.2809	22.737	13.163	0.000	254.498
344.064					

```

=====
=====
Omnibus:                17.219    Durbin-Watson:
2.259
Prob(Omnibus):          0.000    Jarque-Bera (JB):
53.569
Skew:                   0.018    Prob(JB):
2.33e-12
Kurtosis:               5.267    Cond. No.
1.28e+03
=====
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.28e+03. This might indicate that there are strong multicollinearity or other numerical problems.

*#I used linear regression to predict Total Sales based on Price and Quantity.*

*#The model had high accuracy ( $R^2 = 0.834$ ),*

*# showing that these two features strongly influence sales.*

*# I used statsmodels in Python to find this and both variables were statistically significant.*